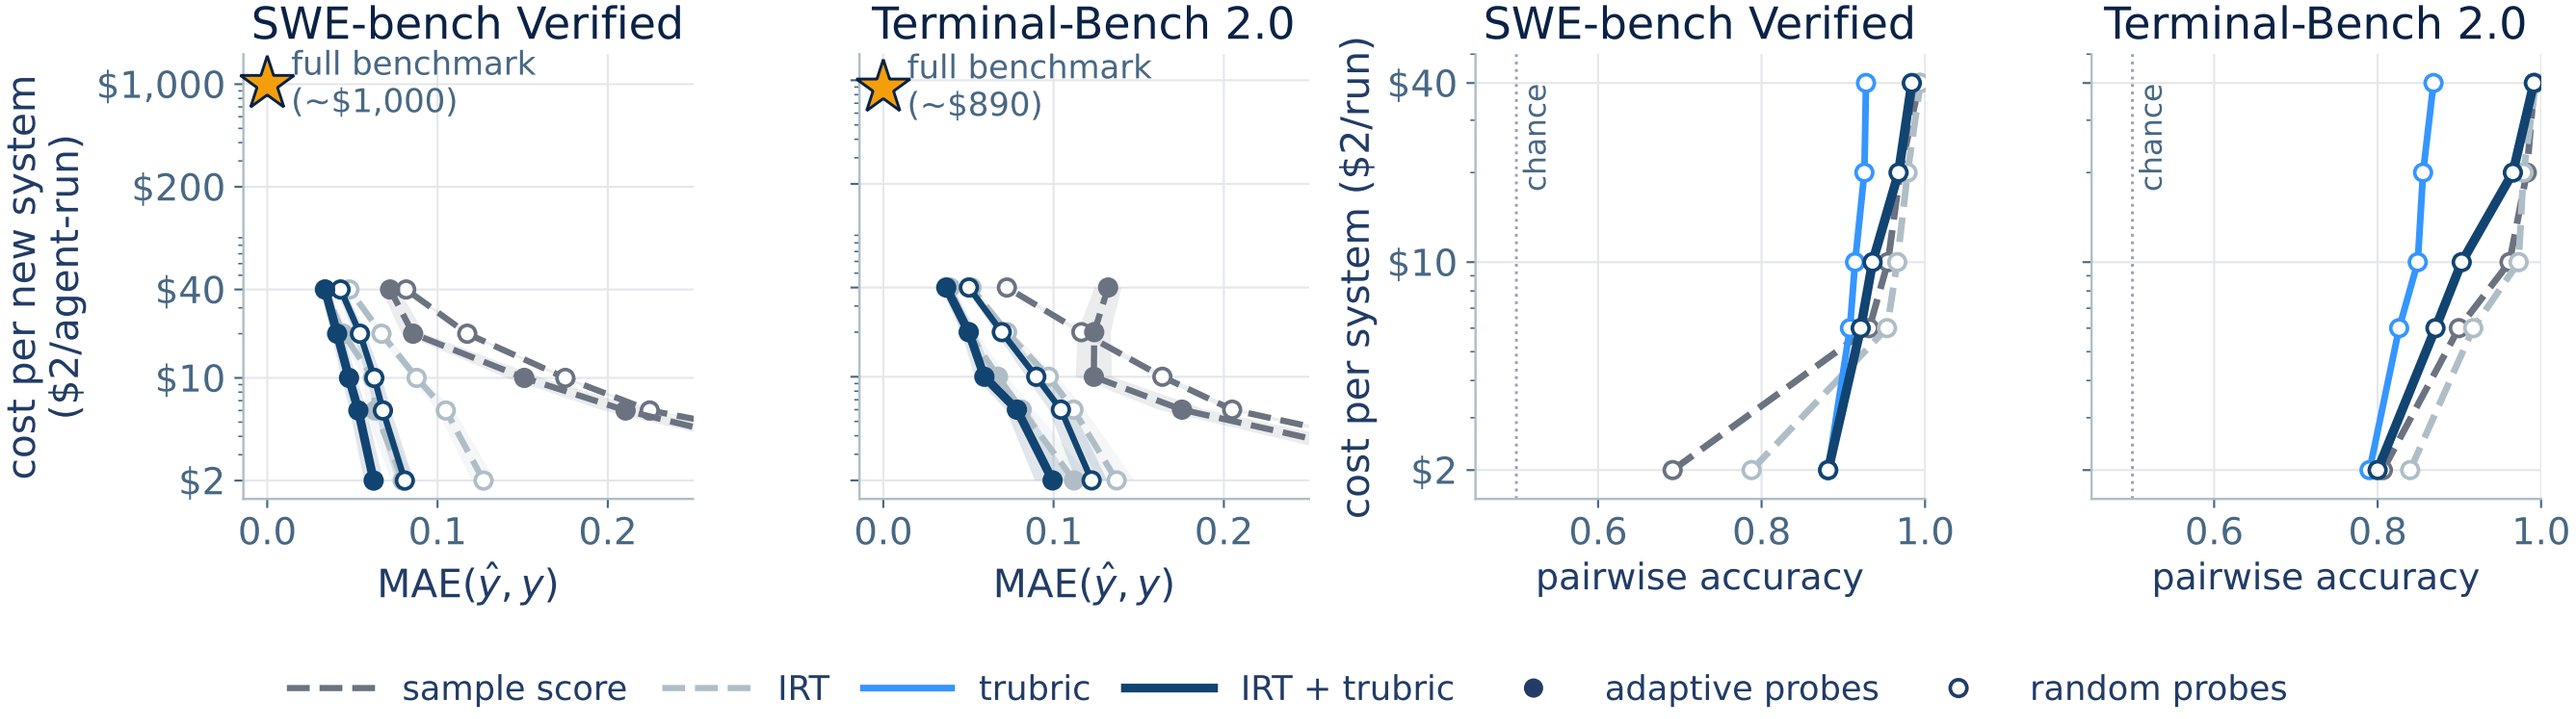

SWE-bench Verified
Terminal-Bench 2.0
SWE-bench Verified
Terminal-Bench 2.0
full benchmark
(~$1,000)
full benchmark
(~$890)
cost per new system
($2/agent-run)
cost per system ($2/run)
$1,000
$200
$40
$10
$2
$40
$10
$2
chance
chance
0.0
0.1
0.2
0.0
0.1
0.2
0.6
0.8
1.0
0.6
0.8
1.0
MAE($\hat{y}$, y)
MAE($\hat{y}$, y)
pairwise accuracy
pairwise accuracy
sample score
IRT
trubric
IRT + trubric
adaptive probes
random probes